

The Google File System

Harris Nagle

CMPT 308 Database Management

April 13th, 2015

Main Ideas About Google File System

1) The Creation of a scalable distributed proprietary file system made in house by Google

- Separate computers are connected together by a network
- Will use large clusters of commodity hardware

2) This system is designed to ...

- Meet the ever-increasing demands of Google
 - Data storage
 - Usage needs
- Be very fault tolerant
- Handle & process multi-TB data
 - Having files be read and written by multiple clients at the same time
 - Manage and organize all files that need to be dealt with on a daily occurrence
- Always favor high bandwidth over low latency!

Implementation

1) Google File System's architecture is composed of the following:

- One master and multiple check servers
- Accessed by multiple clients

2) Master handles all metadata and the operations log into memory

- Makes sure chunk placements and replication decisions using global knowledge.
- Stores three types of metadata:

- 1) File & Chunk namespaces
- 2) Mapping from a file to chunks
- 3) Locations of chunk's replica's

- Files will never be overwritten
- New information will be appended to the file
- Operation Log: Is the logical timeline that defines the sequence of concurrent operations; only persistent record of metadata

3) Divides Files into 64 MB Chunks

- Most file systems chunks sizes vary between 8KB or 1MB

4) Utilizes lazy space allocation

- Chunks are allocated only when necessary

My Take on the Google File System

1) Commodity Hardware

- The use of inexpensive components to power such a widely used search engine on the worlds most visited website.
- Cost Effective
 - Google can continue to add more inexpensive Linux Servers for added storage.

2) Having One Master

- Having only one master simplifies the overall design significantly
- Helps with consistency in respective chunks servers
- One master key is risky, but it's Google, so everything will be fine should things hit the fan.
 - Keeps a history of important and critical metadata changes and replicating each chunk in multiple chunk-servers
 - The master and the chunk-server are both designed to recover itself in seconds; does not matter how either were terminated or exited

Advantages

1) Optimized for large scale data processing

-Larger Chunk Size

- Reduces clients' need to interact with master
- Reduces network overhead
- Reduces the size of the metadata stored in the master
- Metadata is stored in the master's memory
 - It's easy and efficient to periodically scan through its entire background and state.

2) Redundancy

- By default, chunk-servers store three replicas, as discussed earlier, of every chunk
- Provides higher fault tolerance

3) Inexpensive

- Uses commodity hardware
 - Off the shelves Linux servers
- Does not resort to purchasing new hardware or components

Disadvantages

1) Not optimized for smaller work loads

- Can still handle small files and random reads/ writes

2) Single master may cause lesser performance

- May cause a bottleneck
- Everything is reliant on the master
- If the master goes down, so will the entire cluster!

3) Redundancy

- Wastes disk space
- All replicas may not be up to date
- Because files are not replaced when new data is appended, it causes bigger file size

Comparisons

- 1) The large-scale data analysis paper focuses primarily on the map reduce paradigm and then comparing that to the parallel database management systems.
- 2) They both handle large scale data processing abilities
- 3) MapReduce is great because of its simplicity, having only two functions: map and reduce
 - Map: from an input file reads, transforms, then outputs a new key/value pair
 - Reduce: combines the records assigned and then writes the records to an output file
- 4) Parallel DBMS's have all tables partitioned over the nodes in a cluster, and then the system uses an optimizer that translates SQL into execution over multiple nodes